

Chronos: A Content-Addressed Versioned Data Store

A Git-Inspired Approach to Structured Data Versioning Using Merkle Search Trees

Chronos Project Team

Content-Addressed Versioned Data Store Project

github.com/atharvasheersh/Chronos

Abstract-Versioning structured datasets requires durable history, efficient sparse updates, historical reconstruction, and comparison without copying or scanning every record. Chronos is a Python prototype built around an immutable fixed-depth Merkle hash trie. A commit rewrites only affected leaf buckets and the union of their ancestor paths; unchanged subtrees are reused by content hash. Chronos-H adds addressed changesets and selects between log aggregation and hash-pruned tree differencing using a separately calibrated threshold. We evaluate Snapshot, Log-only, forced-Merkle Chronos-M, Chronos-H, and Dolt 2.3.1 on identical deterministic workloads up to 100,000 rows. The audited paper run contains 333 successful and correct trials. Under the CLI-facing protocol, Chronos-H diff latency was 5.8x to 554.2x lower than Dolt and its incremental commit latency was 2.3x to 37.4x lower, while Chronos-H consumed 1.83x to 6.43x more repository storage. Against forced Merkle, the hybrid log path reduced median sparse or hot diff latency by 6.0x to 10.0x. A five-configuration sensitivity study shows that the current $b=8$, $d=4$ trie minimizes storage in the tested set, whereas $b=8$, $d=3$ favors latency. The results support a workload-dependent hybrid design, not universal superiority over a production SQL system.

Index Terms-content-addressed storage, dataset versioning, Merkle trie, copy-on-write, adaptive differencing, SQLite, Dolt

I. INTRODUCTION

Version control for structured data is more than archiving files. A useful system must retain exact versions, apply sparse mutations without rebuilding the complete state, reconstruct historical versions, and identify changes efficiently. Git demonstrates the value of immutable content-addressed objects and parent-linked commits [2], but its tree model is file-oriented. Structured records introduce stable

keys, diverse mutation densities, and comparisons that may span either a few operations or many commits.

Chronos studies whether a deliberately simple fixed-depth Merkle hash trie can combine incremental structural sharing with an operation index. The primary model, Chronos-H, selects a log path for short or sparse history intervals and a Merkle path for intervals whose accumulated operations exceed a frozen calibration boundary.

SQLite supplies an atomic durable object container; it does not define the versioning algorithm.

Dolt is the production comparator because it combines SQL semantics, a Git-style commit graph, and content-addressed Prolly Trees [3], [4]. Chronos is not presented as a feature replacement for Dolt. The comparison instead identifies the performance and storage trade-offs of a smaller instrumentable mechanism under a reproducible CLI-facing protocol.

A. Research Questions

- RQ1: What latency and storage trade-offs does Chronos-H exhibit against Snapshot, Log-only, and Dolt across dataset size, mutation density, and locality?
- RQ2: Does adaptive log/Merkle selection improve diff latency relative to a forced-Merkle ablation without changing correctness?
- RQ3: How sensitive are fixed-trie latency, storage, and examined work to branching factor and depth?

B. Contributions

- A canonical SHA-256-addressed fixed-depth trie with incremental multi-key path rewriting.
- A hybrid diff policy combining addressed changesets with hash-pruned tree comparison.
- An atomic SQLite repository that persists objects, commits, aliases, metrics, and HEAD and verifies hashes after reopen.
- A five-variant benchmark with shared deterministic workloads, correctness oracles,

Dolt integration, raw evidence, and an independently audited summary.

- A controlled branching-factor/depth sensitivity study that prevents treating $b=8$, $d=4$ as universally optimal.

II. LITERATURE REVIEW AND RELATED WORK

Merkle's authenticated tree construction established that a compact root digest can commit to a larger collection [1]. Git applies content-addressed blobs, trees, and commits to software history [2]. Chronos adopts immutable object identity and parent links, but routes structured keys by hash into a persistent trie rather than mirroring a filesystem hierarchy.

Dataset-versioning research exposes a broader storage-retrieval design space. DataHub motivates collaborative dataset history at scale [5]. Decibel integrates branching into a relational storage engine and compares version-first, tuple-first, and hybrid layouts [6]. OrpheusDB bolts versioning onto a conventional DBMS and optimizes partitioning for version retrieval [7]. Bhattacharjee et al. formalize the conflict between storing more materialized state and paying more reconstruction cost [8]. These systems motivate reporting commit, checkout, diff, and storage together rather than selecting a single metric.

Noms represents structured data as a Merkle DAG of immutable chunks and supports efficient diff and synchronization [9]. ForkBase combines content addressing, fork semantics, and duplicate-content detection for forkable applications [10]. Dolt stores table indexes as content-addressed

Prolly Trees: ordered, B-tree-like structures with history-independent chunk boundaries, structural sharing, and native diff [3], [4]. Chronos differs by using fixed-depth hash routing and by exposing an explicit alternative operation-log path. This simplifies instrumentation, but gives up Dolt's ordered range behavior, SQL surface, mature branching, merging, and production engineering.

TABLE I. Literature-review matrix and the gap addressed by Chronos.

Prior work	Core representation	Relationship to Chronos
Git [2]	Object DAG	Identity and commit ancestry; file-oriented
DataHub [5]	Version graph	Collaborative dataset motivation
Decibel [6]	Relational layouts	Native database branching
OrpheusDB [7]	Partitioned versions	Storage/retrieval optimization
Noms [9]	Merkle DAG	Structured content addressing
ForkBase [10]	Forkable CAS	Fork and deduplication semantics
Dolt [3], [4]	Prolly Trees	Production SOTA comparator

III. CHRONOS DESIGN

Chronos-H architecture




Fig. 1. Chronos architecture and research contribution boundary.

A. Addressed Object Model

Trie nodes and changesets are identified by SHA-256 over canonical JSON. A commit hash covers the root, parent, changeset, message, and timestamp. Sequential aliases v1, v2, and so on improve usability but are not content identities. HEAD is a named durable reference to the latest commit.

B. Fixed-Depth Merkle Hash Trie

The evaluated default has branching factor $b=8$ and depth $d=4$. SHA-256(key) contributes three routing bits per level, producing 4,096 possible leaf buckets after twelve bits. Full keys remain in leaves, so routing-prefix collisions do not affect correctness. Expected occupancy under uniform routing is N/b^d , but the sensitivity study treats the geometry as a tunable design choice.

C. Incremental Hashing and Commit

A mutation batch is canonicalized, routed, and grouped by shared prefixes. Each affected leaf is rewritten once, and only the union of affected ancestor paths is rebuilt. Unchanged child hashes

are copied into new internal nodes, providing structural sharing. One SQLite transaction then inserts new objects, the changeset, commit metadata, version alias, measurements, and updated HEAD. Reopen verifies payload hashes and references.

Incremental commit and hybrid diff




Fig. 2. Incremental commit path and adaptive diff decision.

D. Hybrid Diff

The log path walks the ancestor chain and aggregates key operations. The Merkle path recursively compares roots, prunes equal subtree hashes, and examines full leaf entries only below unequal paths. Chronos-H counts accumulated ancestor operations and uses the log at or below 4,096 operations; it uses Merkle above that boundary. Chronos-M forces the Merkle path and therefore serves only as an ablation.

IV. EXPERIMENTAL METHODOLOGY

A. Variants and Workloads

Five implementations receive identical logical states and mutation batches: a complete Snapshot per version, an initial state plus operation-log baseline, forced-Merkle Chronos-M, adaptive Chronos-H, and Dolt 2.3.1 using a keyed SQL table and native

diff. Chronos-M is not a headline product variant; it isolates the selector.

TABLE II. Final evaluation workload matrix.

Scenario	Rows	Commits	Changes	Locality
Small sparse	1,000	10	10	spread
Medium sparse	10,000	10	10	spread
Medium dense	10,000	10	1,000	spread
Large sparse	100,000	10	100	spread
Large hot	100,000	10	100	hot

B. Measurement Contract

Initial import measures creation of the first durable state; incremental commit is the within-trial median across ten durable commits; diff consumes the complete initial-to-final comparison; checkout fully materializes the target version; storage is the sum of regular repository-file bytes. Correctness requires exact initial and final states plus a changed-key oracle. Work examined uses each system's natural unit and is not normalized across keys, operations, and trie-node pairs.

C. Repetition, Calibration, and Integrity

The paper run uses two warmups and seven measured trials with base seed 20260822. Every adapter receives a regenerated workload whose SHA-256 must match the manifest. Calibration

workloads are disjoint from evaluation. Medians, 25th percentiles, and 75th percentiles exclude warmups.

Run 7048cc1674dc4038a8e42c87818b9088 was executed on Windows-11-10.0.26200-SP0 with Python 3.14.3. An independent audit reconstructed the expected trial matrix and every summary statistic.




Fig. 3. Independent threshold calibration; 4,096 is the largest sampled log-favorable point.

The boundary is a frozen policy for this machine, not proof that the true crossover is exactly 4,096. All 333 paper-profile rows completed with status ok and correctness true. The audit retained all 83 Tukey outlier candidates; no observation was manually deleted.

V. RESULTS

A. Final Diff Performance




Fig. 4. Final diff latency from the audited CSV bundle; lower is better.

TABLE III. Chronos-H versus Dolt median diff latency.

Scenario	Path	H ms	Dolt ms	Ratio
Small sparse	log	0.417	188.91	453.3x
Medium sparse	log	0.553	306.47	554.2x
Medium dense	merkle	56.496	326.01	5.8x
Large sparse	log	3.082	180.77	58.7x
Large hot	log	2.031	182.41	89.8x

Chronos-H selected log for the four sparse or hot workloads and Merkle for the 10,000-operation medium-dense workload. Relative to Dolt, its median diff was 5.8x lower in medium-dense and 58.7x to 554.2x lower in the other workloads. These ratios include repeated Dolt CLI and SQL invocation costs and therefore describe the evaluated system protocol, not isolated Prolly Tree algorithm speed.

B. Hybrid Ablation

For log-selected workloads, Chronos-H reduced median diff latency relative to forced-Merkle Chronos-M by 6.0x (small sparse), 8.7x (medium sparse), 8.8x (large sparse), and 10.0x (large hot). In medium-dense, both used Merkle and Chronos-H was within 7.3% of the forced variant. This supports the selector's intended mechanism without claiming that its threshold is hardware-independent.

C. Commit, Checkout, and Storage




Fig. 5. Chronos-H and Dolt operational latency; lower is better.

Chronos-H incremental commit medians were 2.3x to 37.4x lower than Dolt, and checkout medians were 35.8x to 123.4x lower. Initial import was mixed: Chronos-H was 4.9x faster for small sparse and 1.14x faster for medium sparse, approximately tied for medium dense, and 2.73x to 2.79x slower at 100,000 rows. The latency advantage is therefore not uniform.

Chronos-H paid a consistent storage cost. Its median repository footprint was 1.83x, 3.16x, 2.19x, 4.10x, and 6.43x Dolt's for small sparse through large hot, respectively. The 100,000-row Chronos repositories occupied about 34.4 MiB. This reflects inline values, JSON object encoding, and

the absence of compact chunk packing and garbage collection.

D. Trie-Parameter Sensitivity




Fig. 6. Fixed-trie sensitivity; the preferred point depends on the objective.

All 45 sensitivity trials were correct. The b8-d3 configuration delivered the lowest median diff latency (20.15 ms versus 44.43 ms for b8-d4) and compared 84.3% fewer node pairs, but consumed 61.2% more storage and examined three times as many leaf entries. The current b8-d4 default used the least storage (4.27 MiB). The over-deep b8-d5 design was slower for every timed operation and used 44.8% more storage. Thus b8-d4 is a storage-conscious balanced default, not a universal optimum.

VI. DISCUSSION

The results answer RQ1 with a trade-off rather than a single winner. Chronos-H offered low diff, commit, and checkout latency under the benchmark interface, but its import advantage disappeared at 100,000 rows and its repository was always larger than Dolt's. Snapshot and Log-only also remained important: their simpler representations were often

cheaper to import or store, while Snapshot diff scaled with materialized key count.

RQ2 is supported more directly. Because Chronos-H and Chronos-M share the same durable trie, their sparse-workload diff gap isolates adaptive selection more cleanly than an absolute comparison with Dolt. The log path avoided thousands of trie comparisons when only 100 or 1,000 operations separated the endpoints; the dense case switched to Merkle.

RQ3 rejects a universal geometry. Shallow routing lowers internal traversal but creates larger leaf buckets and higher storage; deeper routing lowers occupancy but increases internal-node work. The appropriate point therefore depends on whether latency, storage, or worst-case bucket scanning dominates the target workload.

The Dolt comparison must remain scoped. Dolt performs SQL parsing, schema management, version-control operations, and production durability work that Chronos does not. Its Prolly Trees additionally support ordered access, branches, merges, and remotes. The reported system-level ratios demonstrate the behavior of these implementations through their declared interfaces; they do not establish that fixed tries are intrinsically faster than Prolly Trees.

VII. LIMITATIONS AND THREATS TO VALIDITY

- Runtime and interface asymmetry: Chronos is Python and in-process, while Dolt is a mature Go executable repeatedly invoked through CLI/SQL boundaries.

- Feature asymmetry: Chronos has linear history and one writer; it does not implement SQL, branches, merges, remotes, schema evolution, or production concurrency.
- Synthetic scope: deterministic key/value workloads reach 100,000 rows but do not represent all real schemas, skew, payload sizes, or long histories.
- Single environment: results come from one Windows 11 machine; cache, scheduling, filesystem, and antivirus behavior may affect short operations.
- Calibration boundary: 4,096 is the largest sampled log-favorable point, not a measured universal crossover.
- Memory comparability: Python tracemalloc values and Dolt process memory are not equivalent, so the paper does not present a cross-system memory graph.
- Work-unit mismatch: keys, operations, trie-node pairs, and Dolt-internal work cannot be interpreted as one normalized counter.
- Outlier policy: all 83 Tukey candidates were retained. Medians and interquartile ranges reduce their influence, but seven trials do not establish distributional normality.

VIII. FUTURE WORK

The immediate improvement is workload-aware structure selection. A repository could estimate key count, prefix skew, mutation density, and leaf occupancy and then choose among tested trie geometries instead of fixing b and d globally. The hybrid selector should likewise learn a crossover

from periodic calibration and include history distance, changed-prefix density, and observed I/O cost rather than only operation count.

Storage efficiency can be improved through compact binary node encoding, value chunking, pack files, compression, object reachability analysis, and garbage collection. An ordered Merkle index or content-defined chunk boundaries would add range scans and history-independent layout while retaining the current fixed trie as a controlled baseline.

System scope should expand to named branches, merge commits, conflict reporting, schema-aware records, concurrent readers and writers, and remote synchronization. Evaluation should add public structured datasets, one-million-row and longer-history workloads, Linux replication, and unified external process telemetry for time, CPU, I/O, and peak resident memory.

IX. CONCLUSION

Chronos demonstrates a complete content-addressed versioning path for structured key/value data: canonical immutable objects, incremental copy-on-write Merkle trie updates, atomic SQLite persistence, historical checkout, and adaptive differencing. The audited 333-trial comparison shows that Chronos-H can substantially reduce CLI-facing diff, commit, and checkout latency relative to Dolt 2.3.1, but it also exposes slower large imports, higher storage consumption, a narrower feature set, and runtime/interface asymmetry. The ablation supports hybrid log/Merkle selection, while the sensitivity experiment shows that trie geometry

must be chosen by workload objective. The defensible contribution is therefore an instrumentable hybrid design and reproducible evidence of its trade-offs, not a claim of universal replacement for a production version-controlled SQL database.

ACKNOWLEDGMENT

The authors acknowledge the open-source Dolt project and the researchers whose dataset-versioning systems informed the experimental design.

REFERENCES

- [1] R. C. Merkle, "A digital signature based on a conventional encryption function," in *Advances in Cryptology - CRYPTO '87*, LNCS 293, 1988, pp. 369-378, doi: 10.1007/3-540-48184-2_32.
- [2] S. Chacon and B. Straub, "Git internals - Git objects," in *Pro Git*, 2nd ed. [Online]. Available: <https://git-scm.com/book/en/v2/Git-Internals-Git-Objects>
- [3] DoltHub, "Dolt storage engine: Block store," Dolt Documentation. [Online]. Available: <https://www.dolthub.com/docs/architecture/storage-engine/block-store/>
- [4] DoltHub, "Dolt's storage engine: Prolly Trees and commit graph," Feb. 2024. [Online]. Available: <https://www.dolthub.com/blog/2024-02-29-storage-engine/>
- [5] A. Bhardwaj, S. Bhattacharjee, A. Chavan, A. Deshpande, A. J. Elmore, S. Madden, and A. G. Parameswaran, "DataHub: Collaborative data

science and dataset version management at scale," arXiv:1409.0798, 2014.

manifests; no table or graph uses an earlier validation export.

- [6] M. A. Maddox et al., "Decibel: The relational dataset branching system," *Proc. VLDB Endow.*, vol. 9, no. 9, pp. 624-635, 2016, doi: 10.14778/2947618.2947619.
- [7] S. Huang, L. Xu, J. Liu, A. J. Elmore, and A. Parameswaran, "OrpheusDB: Bolt-on versioning for relational databases," *Proc. VLDB Endow.*, vol. 10, no. 10, pp. 1130-1141, 2017.
- [8] S. Bhattacharjee, A. Chavan, S. Huang, A. Deshpande, and A. Parameswaran, "Principles of dataset versioning: Exploring the recreation/storage tradeoff," *Proc. VLDB Endow.*, vol. 8, no. 12, pp. 1346-1357, 2015.
- [9] Attic Labs, "Noms technical overview," GitHub. [Online]. Available: <https://github.com/attic-labs/noms/blob/master/doc/intro.md>
- [10] S. Wang et al., "ForkBase: An efficient storage engine for blockchain and forkable applications," *Proc. VLDB Endow.*, vol. 11, no. 10, pp. 1137-1150, 2018, doi: 10.14778/3231751.3231762.

REPRODUCIBILITY AND DATA

AVAILABILITY

The repository contains the benchmark harness, evidence audit, trie-sensitivity harness, and paper generator. The final evidence bundle is `output/benchmarks/paper-final-20260824`; the sensitivity bundle is `output/benchmarks/trie-sensitivity-20260824`. Results in this paper are generated programmatically from their summary CSV files and verified against raw results and